

Table 1: One-pass residual readout vs. fairly-conditioned autoregressive loop (accuracy). Readout = small MLP on the frozen final residual (4-seed vote); loop = next-token log-prob over the closed answer set, full context, balanced few-shot. Δ = readout $-$ loop.k, paired McNemar on identical rows ($*p < 0.05$, $**p < 0.01$, $***p < 0.001$). All runs batch 8; Qwen3 0.6B/4B on H200, Qwen3-8B/Mistral-7B/Granite-3.1-8B on B200, DeepSeek-V4-Flash on B300. ^a BoolQ loops: Qwen at pad 8192 with best $k \in \{0, \dots, 64\}$, Mistral/Granite/DSV4 at pad 2048, $k=8$. ^b RuleTaker n2k: readout scored on the loop’s exact 400 rows. ^c ARC-Challenge: full test (1165 rows), $k=8$, pad 2048.

Task	Model	Readout	loop.0	loop.k	Δ	p (McNemar)
BoolQ ^a	Qwen3-0.6B	0.753	0.631	0.715	+0.038***	2.4e−05
	Qwen3-4B	0.862	0.854	0.869	−0.007	0.22
	Qwen3-8B	0.879	0.862	0.886	−0.007	0.24
	Mistral-7B	0.841	0.798	0.852	−0.011	0.21
	Granite-3.1-8B	0.854	0.815	0.864	−0.010	0.64
	DeepSeek-V4-Flash	0.896	0.888	0.906	−0.009	0.39
RuleTaker n2k ^b	Qwen3-0.6B	0.638	0.600	0.545	+0.093**	3.7e−03
	Qwen3-4B	0.738	0.653	0.698	+0.040	0.20
	Qwen3-8B	0.753	0.675	0.730	+0.023	0.45
	Mistral-7B	0.660	0.515	0.575	+0.085*	0.017
	Granite-3.1-8B	0.828	0.558	0.695	+0.133***	1.3e−06
	DeepSeek-V4-Flash	0.763	0.605	0.838	−0.074**	1e−03
ARC-Challenge ^c	Qwen3-0.6B	0.492	0.502	0.597	−0.106***	3.2e−11
	Qwen3-4B	0.837	0.826	0.882	−0.046***	2.9e−06
	Qwen3-8B	0.908	0.902	0.913	−0.005	0.57
	Mistral-7B	0.728	0.750	0.779	−0.051***	9.3e−05
	Granite-3.1-8B	0.708	0.732	0.790	−0.082***	5.8e−10
	DeepSeek-V4-Flash	0.951	0.961	0.959	−0.007	0.23

Table 2: Readout probe-placement sweeps (18 cells, all batch 8): best mid-depth tap vs the final-layer readout, paired McNemar on identical rows. *Layers* = stack depth (28/36/36/32/40/43); taps are residual-layer indices. *Plateau onset* = first swept layer within 1 pt of the final-layer readout — the verdict representation is fully assembled from there on; at DeepSeek-V4 this happens by L22/42 (52% depth) on *all three tasks*, with ARC showing the sharpest snap-in (near-chance at L15 $\rightarrow$ 0.944 at L22). Mid-depth taps win significantly at 9/18 cells and never lose significantly.

Task	Model	Layers	Best tap	Depth	Plateau onset	mlp (best)	mlp (final)	Δ	p
BoolQ	Qwen3-0.6B	28	L18	64%	L18 (67%)	0.761	0.746	+0.010	0.18
	Qwen3-4B	36	L24	67%	L24 (69%)	0.868	0.856	+0.011*	0.038
	Qwen3-8B	36	L30	83%	L24 (69%)	0.889	0.878	+0.011*	0.016
	Mistral-7B	32	L16	50%	L16 (52%)	0.863	0.841	+0.022***	1.1e−04
	Granite-3.1-8B	40	L20	51%	L17 (44%)	0.868	0.854	+0.018**	1.0e−03
	DeepSeek-V4-Flash	43	L36	86%	L22 (52%)	0.900	0.896	+0.005	0.20
RuleTaker n2k	Qwen3-0.6B	28	L18	64%	L13 (48%)	0.688	0.650	+0.036*	0.020
	Qwen3-4B	36	L24	67%	L24 (69%)	0.781	0.744	+0.038**	2.8e−03
	Qwen3-8B	36	L24	67%	L24 (69%)	0.778	0.758	+0.025*	0.040
	Mistral-7B	32	L18	56%	L11 (35%)	0.702	0.643	+0.059***	5.2e−05
	Granite-3.1-8B	40	L20	51%	L20 (51%)	0.816	0.805	+0.011	0.37
	DeepSeek-V4-Flash	43	L29	67%	L22 (52%)	0.776	0.765	+0.021	0.057
ARC-Challenge	Qwen3-0.6B	28	L26	93%	L26 (96%)	0.500	0.492	+0.013	0.082
	Qwen3-4B	36	L29	81%	L24 (69%)	0.856	0.841	+0.022***	2.2e−04
	Qwen3-8B	36	L35	97%	L24 (69%)	0.909	0.909	+0.001	1.00
	Mistral-7B	32	L20	63%	L18 (58%)	0.741	0.728	+0.013	0.50
	Granite-3.1-8B	40	L39	100%	L39 (100%)	0.708	0.708	0.000	—
	DeepSeek-V4-Flash	43	L42	98%	L22 (52%)	0.951	0.951	−0.005	0.29