

The Vector Is the Answer: One-Pass Readouts of Frozen Language-Model Residuals Match the Autoregressive Loop

Hanan Herzog
hanan.herzog@gmail.com

Abstract

On tasks whose answer is representable in a single forward pass, the autoregressive loop is not the only readout of a language model’s computation — and, we show, not a privileged one. We fit small supervised heads to the frozen final-token residual of six open-weight models (Qwen3-0.6B/4B/8B, Mistral-7B, Granite-3.1-8B, DeepSeek-V4-Flash) and run the readout head-to-head against a fairly-conditioned autoregressive loop — full context, balanced few-shot, scored by next-token log-probability over the answer set — matched on information access, scoring protocol, and supervision budget. On BoolQ (full validation), RuleTaker (n2k), and ARC-Challenge (full test), the one-pass readout **matches or beats** the fair loop at nearly every scale: it wins significantly where the model is weakest (RuleTaker +12.0pt on Granite-8B, +5.2pt at 0.6B; BoolQ +3.0pt at 0.6B against a $4\times$ -budget loop), ties at mid scale, and loses decisively only in three regimes — small models on parametric knowledge (ARC), a frontier MoE on serial deduction, and Mistral-7B on BoolQ. The result survives label-shuffle nulls, random-projection selectivity, per-length deconfounds, and paired significance testing on identical rows. Sweeping the readout tap across depth, the verdict representation is fully assembled by mid-depth: the best mid-layer tap beats the final layer significantly at 9 of 18 task $\times$ model cells and never loses. The model, in short, knows more than it can say [15, 21] — on this task class, **the vector already has the answer**. We bound the claim explicitly: it is readout-vs-loop on a frozen backbone, not readout-vs-fine-tuned-specialist, and it does not exhibit a regime where a frozen loop is required.

1 Introduction

Autoregressive decoding is the universal serving interface for language models, but it is only one readout of the computation. Every generated token is produced by applying a fixed, next-token-trained unembedding to a residual vector; the answer a model *would* give is thus a particular linear read of a vector that exists, in full, after a single forward pass — and our own linear probes sit within a few points of the MLP everywhere (§3). This paper asks a sharp, falsifiable question about that vector: **on tasks whose answer can be represented in one forward pass, does a small readout head fit to the frozen residual recover the answer — and does the autoregressive loop add anything such a readout lacks?**

The question is easy to answer sloppily. A probe that beats a zero-shot prompt has shown little: zero-shot understates the loop, and a token-probability decision rule is a known-weak baseline [7]. A probe compared only to the model’s own self-report has shown something else entirely — that beliefs are present, not that a readout can *replace* the loop at producing the right answer. And any probe result without label-shuffle nulls, random-projection selectivity, and significance testing is one confound away from collapse. Our contribution is therefore as much a **discipline of controls** as a result: we run the comparison the probing literature has stopped short of — frozen open weights, a readout fit to gold labels, against a fairly-conditioned loop, matched on supervision budget, information, and scoring.

Concurrent work. The observation that activations carry content the response does not report is now independently published: Hazenoot et al. [18] show a linear probe on frozen activations beats the same model’s own prompted yes/no answer on ESG concept measurement (11/12 comparisons), with the winning layer mid-network. Their loop baseline is zero-shot prompting — which we show understates the loop (§3, Ext 13) — and the comparison carries no paired significance or selectivity controls. We take the convergence as evidence the premise is ripe; the measurement we run here is the part that was missing.

Findings. On BoolQ (full validation, 3270 rows), RuleTaker (n2k pilot), and ARC-Challenge (full test), across six models from 0.6B to a frontier MoE:

1. **Parity or better almost everywhere.** The one-pass readout matches or beats the fair few-shot loop at nearly every scale — winning significantly exactly where the loop is weakest, and tying at mid scale (§3, §4).
2. **The loop wins in three bounded regimes:** small models on parametric knowledge (ARC-Challenge at ≤ 7 B scale), a frontier MoE on formal serial deduction (RuleTaker at DeepSeek-V4), and Mistral-7B on BoolQ. The first two gaps close or invert with scale and task type (§5).
3. **The answer is assembled by mid-depth.** Sweeping the readout tap across the residual stream, the best mid-layer tap beats the final layer significantly at 9/18 cells (6/18 after Benjamini–Hochberg) and never loses — the final layer is a *lower bound* on what the residual readout can do (§4.4).
4. **The margins can be large.** Where the loop’s verbal channel is weakest, the readout’s lead is largest: +12.0pt on RuleTaker with Granite-3.1-8B (readout 0.805 vs loop 0.685, McNemar $p = 6.5\text{e}-12$).

Claim discipline. The readout matches or beats the loop in 15 of 18 cells. The three loop wins are a frontier MoE on serial deduction (RuleTaker at DeepSeek-V4, -3.8pt), Mistral-7B on BoolQ (-2.2pt), and small models on parametric knowledge (ARC, $-10.6 \rightarrow -0.7\text{pt}$ as models grow). The remaining BoolQ cells are statistical ties — four of six sit within $\pm 1.2\text{pt}$, all non-significant. The contribution is the *map* of where each readout wins, not a blanket verdict — and the map is mostly one color.

Why it matters. Later work on residual heads, latent taps, and efficient inference can treat “the answer is already in the vector” as established on a defined task class — not as a serving system we ship here, and not as a claim about all tasks. We state the boundaries explicitly (§5, §6): fine-tuned specialist encoders beat our readout; serial-depth tasks may require the loop; the frozen continuous-latent loop drifts. Within those bounds, the footnote is earned.

Scope. A companion position paper explores a serving architecture of kilobyte-scale readout heads (Paper 2). This paper contains no architecture claims and no production system.

2 Setup and Method

2.1 Two readouts of the same vector

One-pass readout. Let $h_T \in \mathbb{R}^d$ be the final-token residual of a frozen model after one forward pass over the task prompt (question + passage, truncated at `pad_max = 384`). We fit a head $f(h_T) \rightarrow \hat{y}$: either multinomial logistic regression or a one-hidden-layer MLP (width 128, 4-seed vote), both on standardized features. The head sees only task labels; the backbone is never updated.

Fair loop. The autoregressive baseline is *not* greedy decode — it is the strongest *classifier* the loop provides. The model receives the full context (balanced few-shot exemplars, `loop_pad_max` 2048–8192) and is scored by next-token log-probability over the closed answer set (Yes/No, or A–D). We report `loop.0` (zero-shot) and `loop.k` (k -shot); where noted we sweep k up to 64 with a $4\times$ pad budget (Ext 13). *Scope:* we

evaluate the strongest *classifier* the loop offers, not a chain-of-thought loop that generates intermediate tokens before answering. Generated-then-answered chains are a different readout, break the matched closed-set scoring protocol, and — on a serial task such as RuleTaker — are not one-pass-representable by construction; they are out of scope for the matched comparison.

2.2 Matching requirements

A readout-vs-loop comparison is only meaningful if neither side is handicapped. We match on three axes: **information access** (both see the same question and passage; the loop additionally sees k exemplars — an advantage we grant it), **scoring protocol** (both produce a distribution over the same closed answer set; accuracy on identical rows), and **supervision budget** (Ext 13 sweeps the loop’s exemplar budget to $4\times$ the readout’s training budget at 0.6B scale).

2.3 Controls

Every headline number carries four controls:

- **Label-shuffle null:** the head refit on permuted labels must sit at chance, ruling out leakage and spurious separability.
- **Random-projection selectivity:** a head on random Gaussian projections of the residual. If it retains accuracy, the signal is *diffuse* — no privileged subspace — which is what we find (permuted-basis $\approx$ max-basis $\gg$ noise floor) at every scale.
- **Final-token vs full-context:** a mean-pooled full-context readout must underperform the final-token readout, confirming the verdict is *computed into* the final state rather than read off the passage surface.
- **Paired significance:** McNemar’s exact test and paired bootstrap on per-row predictions from identical val rows, not pooled-accuracy eyeballing.
- **Per-length deconfound:** headline cells are re-scored within question- and passage-length terciles; accuracy is flat across length (0.89–0.90 on BoolQ at DeepSeek-V4), so the readout’s edge is not an artifact of the loop truncating long rows.

2.4 Models and artifacts

Six open-weight models spanning three families and a width ladder: Qwen3-0.6B (28×1024), Qwen3-4B (36×2560), Qwen3-8B (36×4096), Mistral-7B (32×4096), Granite-3.1-8B (40×4096), DeepSeek-V4-Flash (43×4096 , FP8 MoE). All runs are batch 8 on model-pinned GPUs (0.6B/4B H200, 8B/Mistral/Granite B200, DSV4 B300). Trained heads are persisted per layer as versioned artifacts (~ 0.5 – 8.4 MB each); per-row predictions persist for paired tests. Run manifests embed the batch tag in cache keys after the Ext 15 cache-identity incident.

3 BoolQ: one-pass readout vs the fair loop

BoolQ [9] is binary reading comprehension with the gold answer grounded in a passage — the cleanest case of a one-pass-representable task. We load the standard Hugging Face release (google/boolq) and use the full validation split (9427 train / 3270 val).




Figure 1: BoolQ: readout vs loop across six models (full val).




Figure 2: Supervision budget (Ext 13). The loop’s k -curve plateaus by $k=32$ and never passes the readout.

3.1 Headline

Across the Qwen and Granite scales, the one-pass readout **wins at 0.6B, ties at 4B/8B, and trails the few-shot loop by ~ 1 pt (ns) at Granite-8B and the frontier MoE** (Table 1). The one exception is Mistral-7B, where the loop wins by 2.2pt ($p = 2.6e-03$) — the single significant loop win on BoolQ, which we flag as a single-model exception (§6).

3.2 Supervision asymmetry is bounded (Ext 13)

The natural objection: the head sees thousands of labels, the loop sees k exemplars. We gave the loop its budget back — k swept to 64 balanced exemplars at $4 \times$ pad (8192) — at 0.6B/4B/8B. The loop’s k -curve **plateaus by $k=32$** (loop.64 – loop.32 = -0.0003 , $p = 1.00$ at 0.6B) and **never passes the readout where the readout was winning**: at 0.6B the readout’s +3.0pt over the best loop stays significant ($p = 3.2e-05$); at 4B/8B all cells are exact ties. Supervision starvation is not the explanation (Fig. 2).

3.3 Readings

1. **Parity or better at five of six models; the edge is largest where the model is weakest.** The readout’s only significant win is at 0.6B — the loop’s verbal channel is the bottleneck there, and the readout channel saturates before it. Mistral-7B is the exception (loop +2.2pt).
2. **The signal is distributed, not localized.** Random-projection heads retain accuracy (perm $\approx$ max $\gg$ noise) at every scale; no privileged subspace carries the verdict (Table 4).
3. **The verdict concentrates in the final token.** Mean-pooled full-context readouts underperform the final-token readout everywhere (full-context $\ll$ final-token) — the answer is computed into the final state, not read off the passage.
4. **The readout is essentially linear.** A multinomial logistic readout tracks the MLP within a few points at every scale (BoolQ 0.6B linear.max 0.750 vs MLP 0.747; 8B 0.874 vs 0.881), so the verdict is a near-linear read of the residual, not a fragile nonlinear probe. At 0.6B, where the residual does not yet hold the answer (ARC $\approx$ 0.49, near chance), no head can invent it — the readout reads what is there.

Context length. The readout reads at pad_max = 384 against the loop’s 2048–8192. Because the BoolQ gold answer is grounded in the passage, left-truncation could in principle drop long passages asymmetrically; we do not log per-row truncation rates and flag this as a limitation (§7), noting the per-length deconfound above shows no length-dependent accuracy collapse.

3.4 FLOPs, precisely

Readout and *scoring* loop are both a single context forward pass — FLOP-equal up to context length. The readout’s structural lever is against a *decoding* loop — and here the non-self-termination is **measured**, not assumed: a greedy decode from the “Answer:” prompt never stops (25/25 hit the 300-token cap at 0.6B; 7/8 hit the 1000-token cap at DeepSeek-V4). The parametric win is therefore real and unconditional: **a head instead of a decode loop — no new tokens, no KV-cache.** (We quantify the economics and their boundary against early-exit serving in Table 3; the decode row is measured.)

4 RuleTaker and ARC: serial depth and parametric knowledge

BoolQ is passage-grounded. We test the two ways a task can be harder for a one-pass readout: **serial deduction depth** (RuleTaker) and **parametric knowledge with no passage** (ARC-Challenge).

4.1 RuleTaker n2k (pilot): matched loop rows, per-depth strata

RuleTaker [11] is closed-world rule reasoning with labeled deduction depth (0/1/2/3/5 + NatLang), loaded via Hugging Face `tasksource/rulemaker`. Protocol: 2000 train / 1000 val (n2k pilot subsample of the full release), loop on the same 1000 rows, $k \in \{0, 8\}$, batch 8. On Qwen, the matched one-pass MLP **beats the fair 8-shot loop at every scale** (+3 to +5pt; significant at 0.6B and 4B, $p = 3.8\text{e-}03$ and $5.3\text{e-}03$; 8B is a near-tie $p = 0.056$). The cross-family models sharpen the picture: Mistral-7B +7.2pt ($p = 8.4\text{e-}04$), **Granite-3.1-8B +12.0pt** ($p = 6.5\text{e-}12$) — **the largest readout margin in the campaign.** The one clean loop win is DeepSeek-V4 (−3.8pt, $p = 8.5\text{e-}03$), driven by the shallow-depth strata.

Per-depth (same rows): no sharp “only the loop works past depth D ” cliff. Accuracy degrades gradually with depth; $d=5$ remains well above chance. The DSV4 loop win concentrates at shallow depths ($d0/d1$, +10–15pt); at 0.6B the loop *collapses* with depth ($d5$ loop.8 0.392 $\ll$ readout 0.544). Thin bins ($n = 21\text{--}51$) are reported with n and not over-read (Fig. 3).




Figure 3: RuleTaker overall + per-depth strata, readout vs loop.

4.2 ARC-Challenge: the loop’s best case

ARC-Challenge [10] is the hard partition of the AI2 Reasoning Challenge (not Chollet’s Abstraction & Reasoning Corpus): grade-school science multiple-choice questions that retrieval and co-occurrence baselines miss. We load `allenai/ai2_arc` config ARC-Challenge and keep 4-choice A–D items only ($\sim 99\%$ of the set). Answers are parametric science knowledge with no passage — they must come from the weights, not the context. Here the loop wins at small scale: -10.6pt at **0.6B**, -4.6pt at **4B** (both $p < 1e-05$), narrowing to -0.5pt at 8B and -0.7pt at DSV4 (both ns). The cross-family models replicate the pattern: Mistral-7B -5.1pt ($p = 9.3e-05$), Granite-8B -8.2pt ($p = 5.8e-10$). The gap **closes monotonically with scale** — few-shot exemplars teach a task format the residual head must otherwise learn from labels, and the advantage shrinks as the parametric answer saturates the residual.

4.3 The task-type regularity

Across 18 task×model cells, a single regularity organizes everything (Fig. 4): **the loop wins in three regimes** — small models on parametric knowledge (ARC), a frontier MoE on formal serial deduction (RuleTaker at DSV4), and Mistral-7B on BoolQ. Everywhere else the one-pass readout wins or ties. The ARC crossover tracks *capability*, not family (Granite/Mistral sit where Qwen did at equal accuracy); the RuleTaker crossover tracks *serial structure at frontier scale*; the Mistral BoolQ cell is a single-model exception we flag



Figure 4: The task-type regularity. readout — best loop (paired, same rows) per task×model cell, with McNemar significance. Green everywhere except three red cells — small models on parametric ARC, the frontier MoE on serial RuleTaker, and Mistral-7B on BoolQ.




Figure 5: Three-task head-to-head (BoolQ · RuleTaker · ARC) across models.

rather than explain (§6).

4.4 Readout placement: the final layer is a lower bound

Every headline number above taps the final residual layer. We swept the tap across depth (10 layers, paired McNemar vs the final-layer readout, trained heads persisted at every tap). The curve is the same everywhere: monotone rise, mid-depth plateau, slight terminal decline (Table 2, Fig. 6).

- **The final layer never wins.** Mid-depth taps beat the final layer significantly at **9/18 cells and never lose significantly** — the headline numbers were a lower bound on the readout, not its ceiling. Across the 18 cells, 6 survive a Benjamini–Hochberg FDR correction at $q < 0.05$ (the three borderline cells — BoolQ 4B, RuleTaker 0.6B, RuleTaker 8B — drop out); the “never loses” direction is unchanged.
- **The optimum is task-stable, not scale-driven.** BoolQ/RuleTaker peak at 50–86% depth at every scale; RuleTaker’s optimum sits at $\sim 2/3$ depth from 0.6B through DSV4. ARC drifts to the final layer at $\geq 8B$ — once the parametric answer saturates the residual, late-layer specialization stops costing anything.
- **Placement converts parity into wins.** Where the loop catches the final-layer readout (BoolQ 8B: L35 ties loop), the best mid-layer tap re-establishes the lead (L30 +1.1pt, $p = 0.016$; on Granite, L20 +1.8pt, $p = 1e-03$).

Readout placement sweeps (all batch 8) — mid-depth taps significantly beat the final layer at 9/18 cells, never lose significantly; shaded band = plateau onset (first layer within 1pt of the final-layer readout)




Figure 6: Readout probe-placement sweeps (18 cells). One-pass readout accuracy vs tapped residual layer; loop arms as horizontal references; paired McNemar annotation per panel. Mid-depth taps win significantly at 9/18 cells and never lose.

- **Wide-shallow plateaus earlier.** Mistral (32×4096) onsets at 35–58% depth vs Qwen3-8B’s uniform 69% at the same width — at fixed width, the shallower stack assembles the verdict sooner. Granite (40×4096 , deep) sits between the two.

Causal hedge. The monotone-rise $\rightarrow$ plateau $\rightarrow$ terminal-decline shape is *consistent with* late-layer specialization toward next-token surface form, but we do not demonstrate it: we never ablated deep layers against generation, and never probed next-token structure per layer. The direct test (a per-layer next-token probe) is future work; “the verdict readout needs less than full depth” and “the final third is load-bearing for decoding” are different claims, and only the first is on the table here.

5 The boundary: where the loop could win, and why it doesn’t cleanly here

A single forward pass is a bounded-depth computation (TC^0 in the circuit- complexity framing); the loop’s advantage should appear exactly where serial depth is irreducible. We hunted that regime and found it only weakly exhibitable on frozen base weights:

- **RuleTaker per-depth** shows a gentle slope, not a regime switch — the loop edges the readout by a few points at $d=5$ (noisy thin bins), and the one decisive loop win (DSV4) concentrates at *shallow* depths, where the frontier model’s in-context deduction is strong.
- **ARC** shows the loop’s real advantage is *parametric retrieval under a learned task format*, not serial compute — and it closes with scale.
- **Synthetic mechanism** (multi-hop transitivity with shuffled facts; Fig. 7): the relation is linearly decodable from the frozen residual (0.93–0.97) where the zero-shot decoder sits at chance; a 2-demo budget surfaces the loop’s inference but never exceeds the one-pass MLP. The measured limit is **order-locked binding**, not depth.

The relation is in the residual; the loop never exceeds a one-pass readout




Figure 7: Mechanism (synthetic). Multi-hop transitivity: one-pass readout vs decoder loop across Qwen3 scale. The loop never exceeds a one-pass readout.

The remaining “loop wins” cell — a TC^0 -hard task on which a *frozen* loop decisively beats one-pass — is largely unoccupied in our sweep; the counter-position (COCONUT, looped-latent) buys serial depth by *training*, which is out of scope here. Because RuleTaker is an n2k pilot with thin per-depth bins ($n = 21\text{--}51$), we treat the serial-depth boundary as *exploratory*: it indicates where a TC^0 gap would appear but is underpowered to fix its location.

6 Discussion

The thesis, in its conditional form. On closed-set, one-pass-representable tasks, a small head on the frozen residual matches a fair autoregressive classifier — so the answer is already in the vector. This is a measurement, not a metaphysics: we do not claim the loop never computes, and we identify the three regimes where it does (§5).

The model knows more than it can say. The readout’s largest margins are exactly where the loop is weakest (Granite RuleTaker +12.0pt; 0.6B everywhere) — the readout channel saturates before the verbal channel. This is the internal-knowledge/external-behavior disconnect Orgad et al. [21] document and Gekhman et al. [15] formalize as *hidden knowledge*; we show it has *operational* content: a producer readout on the prompt residual is a drop-in replacement for the classifier the loop provides. That the readout is near-linear (§3) sharpens the point: the answer is a near-linear read of a vector the model already computes.

The Mistral BoolQ exception. Mistral-7B is the one model where the loop significantly wins on BoolQ (−2.2pt, $p = 2.6e-03$). We report it plainly rather than explain it away: it is a single-model, single-task cell, and Mistral’s wide-shallow stack (32×4096) plateaus earliest of the campaign (§4.4), so its final-layer residual is the weakest readout substrate relative to its own loop. One cell does not move the map, but it is a real boundary, not noise.

Separability of readout from generative interface. The unembedding is one readout of the residual — a fixed, next-token-trained one. That a supervised head matches it is the default expectation once stated that way; the empirical content is *where* it fails (ARC at scale, DSV4 deduction) and that the failure modes are orderly.

Specialist frontier (acknowledged). Fine-tuned encoder classifiers (RoBERTa-large ~ 0.87 , DeBERTa-v3 ~ 0.88 , T5-11B 0.91 on BoolQ) **beat** our 8B one-pass readout (0.878). Our claim is readout-vs-loop on a frozen backbone — the comparison a serving stack faces when the weights are already loaded — not

readout-vs-fine-tuned-specialist. Where a specialist can be trained and shipped, it wins on raw accuracy; the readout’s economics are marginal-cost (the forward pass already happened).

What we are not claiming. A frozen continuous latent loop works (it drifts — Exp 7); the loop never computes; we beat fine-tuned specialists; we shipped residual-first inference. Paper 2 explores the serving architecture as a proposal.

7 Limitations

- **RuleTaker is an n2k pilot** (2000/1000, natural depth mix, $d3$ fat), not the full test; loop rows are 400.
- **Paired tests cover all 18 placement cells and the BoolQ budget arm**, but the per-depth loop strata are thin ($n = 21\text{--}51$), and 3 of the 9 significant placement wins do not survive Benjamini–Hochberg.
- **Truncation is not logged per row**: the readout reads at `pad_max = 384` (left-truncated) against the loop’s 2048–8192, and we do not record the fraction of rows truncated, so a length-asymmetric truncation cost on long-passage rows cannot be ruled out from the logged data (the per-length deconfound shows no accuracy collapse, but it is a coarse check).
- **Supervision asymmetry is bounded, not eliminated**: the head still sees more labels than the loop’s exemplars at scale parity (Ext 13 closes it at 0.6B; Arm C — a fine-tuned decoder — is future work).
- **Reproducibility**: per-layer heads are persisted for all sweep runs; the two oldest canonical headline runs (Ext 8/9) predate head persistence. The FLOP-fraction figure is order-of-magnitude (active-param counts for the MoE are unmeasured — Table 3).
- **Scope**: in-knowledge, closed-set tasks only. No claim transfers to open-ended generation.

8 Related Work

Accessing what a frozen model knows is established from several directions, none of which runs our comparison. *Behavioral* work [24] shows CoT rationales can be unfaithful — text-only access, no representation. *Latent-belief* work reads the representation directly: CCS [4], pre-CoT belief probes [12], belief-formation timing [2], and the correctness/self-knowledge cluster [5, 15, 19, 21]. Orgad et al. detect errors (and, in their §6, re-rank 30 samples); Gekhman et al. formalize *hidden knowledge* and re-rank up to 1k samples with a verifier probe on $h(q, a)$ — both score *candidates the model already emitted*, not a producer head on the prompt residual alone, and neither holds that head to a fairly-conditioned loop under matched supervision. *Decoding-suboptimality* work [1, 3, 7, 8] and, closest, Hazenoot et al. [18] show hidden-state readouts beat the token-probability decision rule or zero-shot prompting — baselines their own analyses show are weak. *Probe-for-verification* work [20] uses internal states to improve generation, not replace it.

The empty cell — frozen open weights, a producer readout on $h(q)$ fit to gold labels, run head-to-head against a fairly-conditioned loop, matched on supervision budget, information, and scoring, under format/shuffle/randproj controls — is the one this paper occupies. Turpin saw behavior, Cox saw belief, Orgad detected errors and re-ranked samples, Gekhman formalized hidden knowledge with a verifier on $h(q, a)$, Cho beat a baseline they showed was weak, Hazenoot beat zero-shot prompting, ReProbe verified steps to generate better; we compare a one-pass producer and the loop at producing the right answer, because the weights are open and frozen.

Delimiting counter-position. COCONUT [17] and looped-latent transformers [22] *train* a continuous latent loop to buy serial depth — they set the boundary we test without training. Early-exit serving [6, 14, 16, 23] operationalizes placement redundancy into a serving optimization; it shares our premise and is prior

Table 1: One-pass residual readout vs. fair AR loop (accuracy). Readout: MLP on frozen final residual (per-seed mean). Loop: closed-set next-token log-prob, full context, balanced few-shot. loop.0 / loop.8 = zero- / 8-shot; *best loop* = stronger arm (k in parentheses; RT 0.6B and ARC DSV4 prefer zero-shot). Δ = readout — best loop on identical rows ($\Delta > 0$ = readout wins; best-of makes Δ conservative). Paired McNemar (* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$). ^aBoolQ val 3270; Qwen $k \leq 64$ pad 8192, else $k=8$ pad 2048. ^bRuleTaker n2k val 1000. ^cARC-Challenge test 1165, $k=8$.

Task	Model	Readout	loop.0	loop.8	best loop	Δ (readout — loop)	p (McNemar)
BoolQ ^a	Qwen3-0.6B	0.746	0.631	0.645	0.715 (32)	+0.030 ***	3.2e−05
	Qwen3-4B	0.856	0.855	0.865	0.867 (64)	−0.012	0.094
	Qwen3-8B	0.878	0.862	0.876	0.886 (16)	−0.008	0.243
	Mistral-7B	0.841	0.824	0.863	0.863 (8)	−0.022**	0.003
	Granite-3.1-8B	0.854	0.800	0.865	0.865 (8)	−0.011	0.076
	DeepSeek-V4-Flash	0.896	0.888	0.906	0.906 (8)	−0.009	0.078
RuleTaker n2k ^b	Qwen3-0.6B	0.650	0.598	0.587	0.598 (zero)	+0.052 **	0.004
	Qwen3-4B	0.748	0.679	0.703	0.703 (8)	+0.045 **	0.005
	Qwen3-8B	0.761	0.666	0.729	0.729 (8)	+0.032	0.056
	Mistral-7B	0.643	0.517	0.571	0.571 (8)	+0.072 ***	8.4e−04
	Granite-3.1-8B	0.805	0.554	0.685	0.685 (8)	+0.120 ***	6.5e−12
	DeepSeek-V4-Flash	0.765	0.603	0.803	0.803 (8)	−0.038**	0.009
ARC-Challenge ^c	Qwen3-0.6B	0.492	0.502	0.597	0.597 (8)	−0.106***	3.2e−11
	Qwen3-4B	0.841	0.826	0.882	0.882 (8)	−0.041***	2.9e−06
	Qwen3-8B	0.909	0.902	0.913	0.913 (8)	−0.004	0.567
	Mistral-7B	0.728	0.750	0.779	0.779 (8)	−0.051***	8.3e−05
	Granite-3.1-8B	0.708	0.732	0.790	0.790 (8)	−0.082***	5.8e−10
	DeepSeek-V4-Flash	0.951	0.961	0.959	0.961 (zero)	−0.010	0.064

art for any “skip the tail” claim, which we do not make. Trajectory-dynamics probes [13] read layer-wise displacement across all tokens for OOD monitoring — a richer readout than our static one, aimed at a different goal.

9 Conclusion

On closed-set, one-pass-representable tasks, a small readout head fit to the frozen residual recovers the answer without generating tokens — and matches the strongest classifier the autoregressive loop provides across the space we measured. Three loop wins remain: a frontier MoE on serial deduction (RuleTaker at DeepSeek-V4), Mistral-7B on BoolQ, and a small-model parametric gap (ARC) that closes with scale; elsewhere the loop’s edges are statistical ties (four of six BoolQ cells within ± 1.2 pt). The claim is therefore bounded but strong: on this task class, the vector already has the answer — and we map the boundaries where the loop still earns its keep. The readout is near-linear — the answer is a near-linear read of a vector the model already computes. **The vector already has the answer.** The construction of serving machinery on that substrate is companion work; here, we measured the substrate.

Acknowledgments

Large language models assisted with running experiments and drafting this manuscript: Grok 4.5 (xAI), DeepSeek-V4-Flash, and Kimi K3. The author designed the study, reviewed all results and text, and takes full responsibility for the claims.

Table 2: Readout probe-placement sweeps (18 cells, all batch 8): best mid-depth tap vs the final-layer readout, paired McNemar on identical rows. *Layers* = stack depth (28/36/36/32/40/43); taps are residual-layer indices. *Plateau onset* = first swept layer within 1 pt of the final-layer readout — the verdict representation is fully assembled from there on; at DeepSeek-V4 this happens by L22/42 (52% depth) on *all three tasks*, with ARC showing the sharpest snap-in (near-chance at L15 $\rightarrow$ 0.944 at L22). Mid-depth taps win significantly at 9/18 cells and never lose significantly.

Task	Model	#	Tap	Dep.	Plateau	best	final	Δ	p
BoolQ	Qwen3-0.6B	28	L18	64%	L18 (67%)	0.761	0.746	+0.010	0.18
	Qwen3-4B	36	L24	67%	L24 (69%)	0.868	0.856	+0.011*	0.038
	Qwen3-8B	36	L30	83%	L24 (69%)	0.889	0.878	+0.011*	0.016
	Mistral-7B	32	L16	50%	L16 (52%)	0.863	0.841	+0.022***	1.1e−04
	Granite-3.1-8B	40	L20	51%	L17 (44%)	0.868	0.854	+0.018**	1.0e−03
	DeepSeek-V4-Flash	43	L36	86%	L22 (52%)	0.900	0.896	+0.005	0.20
RuleTaker n2k	Qwen3-0.6B	28	L18	64%	L13 (48%)	0.688	0.650	+0.036*	0.020
	Qwen3-4B	36	L24	67%	L24 (69%)	0.781	0.744	+0.038**	2.8e−03
	Qwen3-8B	36	L24	67%	L24 (69%)	0.778	0.758	+0.025*	0.040
	Mistral-7B	32	L18	56%	L11 (35%)	0.702	0.643	+0.059***	5.2e−05
	Granite-3.1-8B	40	L20	51%	L20 (51%)	0.816	0.805	+0.011	0.37
	DeepSeek-V4-Flash	43	L29	67%	L22 (52%)	0.776	0.765	+0.021	0.057
ARC-Challenge	Qwen3-0.6B	28	L26	93%	L26 (96%)	0.500	0.492	+0.013	0.082
	Qwen3-4B	36	L29	81%	L24 (69%)	0.856	0.841	+0.022***	2.2e−04
	Qwen3-8B	36	L35	97%	L24 (69%)	0.909	0.909	+0.001	1.00
	Mistral-7B	32	L20	63%	L18 (58%)	0.741	0.728	+0.013	0.50
	Granite-3.1-8B	40	L39	100%	L39 (100%)	0.708	0.708	0.000	—
	DeepSeek-V4-Flash	43	L42	98%	L22 (52%)	0.951	0.951	−0.005	0.29

Table 3: What each access to the frozen residual costs (economics orientation, not a benchmark). Status stamps every row: “measured” = our run protocol; “projected” = forward arithmetic we did not run. Readout & scoring-loop are the same single pass at equal context — the readout’s structural gap is **no generation**, not layer-skipping. Early-exit serving shares our placement premise and is prior art for any “skip the tail” claim.

Access	Cost shape	Generates?	Status
readout (final tap)	1 fwd , short ctx	no	measured
scoring loop “.8”	1 fwd, long ctx	no	measured
early-exit decode	1 fwd, tail truncated	yes	projected / prior art
decode / output	K fwds (per token)	yes	measured (no self-stop)

Note. The decode row *is* measured: greedy decode from the Answer: prompt never self-terminates (25/25 hit the 300-token cap at 0.6B; 7/8 hit the 1000-token cap at DeepSeek-V4).

Table 4: Controls, all 18 cells. Every headline number clears all four controls. *label-shuffle*: the head refit on permuted labels sits at chance. *random-projection*: a head on random Gaussian projections retains accuracy with permuted-basis $\approx$ max-basis $\gg$ noise floor — the signal is *diffuse*. *full-context*: a mean-pooled full-context readout underperforms the final-token readout — the verdict is computed into the final state, not read off the passage. Balanced chance is 0.50 (BoolQ, RuleTaker) and 0.25 (4-way ARC); the *majority-class* baseline is $\approx$ 0.62 (BoolQ), 0.50 (RuleTaker), 0.27 (ARC).

Task	Model	shuffle	rp max	rp perm	rp noise	full-ctx	final
BoolQ	Qwen3-0.6B	0.533	0.734	0.725	0.601	0.667	0.746
	Qwen3-4B	0.542	0.861	0.859	0.604	0.717	0.856
	Qwen3-8B	0.528	0.874	0.874	0.607	0.739	0.878
	Mistral-7B	0.537	0.810	0.803	0.607	0.685	0.841
	Granite-3.1-8B	0.532	0.848	0.849	0.607	0.679	0.854
	DeepSeek-V4	0.522	0.889	0.891	0.602	0.678	0.896
RuleTaker	Qwen3-0.6B	0.511	0.635	0.635	0.489	0.517	0.650
	Qwen3-4B	0.474	0.718	0.726	0.505	0.586	0.744
	Qwen3-8B	0.488	0.736	0.745	0.496	0.575	0.758
	Mistral-7B	0.495	0.645	0.634	0.496	0.503	0.643
	Granite-3.1-8B	0.491	0.756	0.757	0.496	0.545	0.805
	DeepSeek-V4	0.500	0.726	0.716	0.506	0.518	0.765
ARC	Qwen3-0.6B	0.278	0.470	0.470	0.254	0.278	0.492
	Qwen3-4B	0.269	0.802	0.803	0.255	0.349	0.841
	Qwen3-8B	0.287	0.869	0.867	0.251	0.399	0.909
	Mistral-7B	0.259	0.679	0.678	0.251	0.277	0.728
	Granite-3.1-8B	0.284	0.598	0.588	0.251	0.279	0.708
	DeepSeek-V4	0.269	0.942	0.940	0.236	0.439	0.951

References

- [1] Momin Abbas, Yi Zhou, Parikshit Ram, Nathalie Baracaldo, Horst Samulowitz, Theodoros Salonidis, and Tianyi Chen. Enhancing in-context learning via linear probe calibration. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2024.
- [2] Siddharth Boppana, Annabel Ma, Max Loeffler, Raphael Sarfati, Eric Bigelow, Atticus Geiger, Owen Lewis, and Jack Merullo. Reasoning theater: Disentangling model beliefs from chain-of-thought. *arXiv preprint arXiv:2603.05488*, 2026.
- [3] Marcus Buckmann and Edward Hill. Logistic regression makes small LLMs strong and explainable “tens-of-shot” classifiers. *arXiv preprint arXiv:2408.03414*, 2024.
- [4] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023.
- [5] Iván Vicente Moreno Cencerrado, Arnau Padrés Masdemont, Anton Gonzalvez Hawthorne, David Demitri Africa, and Lorenzo Pacchiardi. No answer needed: Predicting LLM answer accuracy from question-only linear probes. In *ICLR Workshop on Principled Design for Trustworthy AI*, 2026.
- [6] Yanxi Chen, Xuchen Pan, Yaliang Li, Bolin Ding, and Jingren Zhou. EE-LLM: Large-scale training and inference of early-exit large language models with 3D parallelism. In *International Conference on Machine Learning (ICML)*, 2024.
- [7] Hakaze Cho, Yoshihiro Sakai, Mariko Kato, Kenshiro Tanaka, Akira Ishii, and Naoya Inoue. Token-based decision criteria are suboptimal in in-context learning. In *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [8] Hyunsoo Cho, Hyuhng Joon Kim, Junyeob Kim, Sang-Woo Lee, Sang-goo Lee, Kang Min Yoo, and Taeuk Kim. Prompt-augmented linear probing: Scaling beyond the limit of few-shot in-context learners. In *AAAI Conference on Artificial Intelligence*, 2023.
- [9] Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. BoolQ: Exploring the surprising difficulty of natural yes/no questions. In *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2019.
- [10] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try ARC, the AI2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.
- [11] Peter Clark, Oyvind Tafjord, and Kyle Richardson. Transformers as soft reasoners over language. In *International Joint Conference on Artificial Intelligence (IJCAI)*, 2020.
- [12] Kyle Cox, Darius Kianersi, and Adrià Garriga-Alonso. Post-hoc reasoning in chain of thought: Decoding and steering pre-committed answers. *arXiv preprint arXiv:2603.01437*, 2026.
- [13] Hamed Damirchi, Ignacio Meza De la Jara, Ehsan Abbasnejad, Afshar Shamsi, Zhen Zhang, and Javen Shi. Truth as a trajectory: What internal representations reveal about large language model reasoning. *arXiv preprint arXiv:2603.01326*, 2026.
- [14] Mostafa Elhoushi, Akshat Shrivastava, Diana Liskovich, Basil Hosmer, Bram Wasti, Liangzhen Lai, Anas Mahmoud, Bilge Acun, Saurabh Agarwal, Ahmed Roman, Ahmed A Aly, Beidi Chen, and Carole-Jean Wu. LayerSkip: Enabling early exit inference and self-speculative decoding. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [15] Zorik Gekhman, Eyal Ben David, Hadas Orgad, Eran Ofek, Yonatan Belinkov, Idan Szpektor, Jonathan Herzig, and Roi Reichart. Inside-out: Hidden factual knowledge in LLMs. In *Conference on Language Modeling (COLM)*, 2025.

- [16] Andrey Gromov, Kushal Tirumala, Hassan Shapourian, Paolo Glorioso, and Daniel A. Roberts. The unreasonable ineffectiveness of the deeper layers. In *International Conference on Learning Representations (ICLR)*, 2025.
- [17] Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space. In *Conference on Language Modeling (COLM)*, 2025.
- [18] Luc Hazenoot, Zhaochun Ren, and Amirhossein Zohrehvand. Measuring concept content in text from LLM activations: ESG evidence from concept vectors and linear probes. *arXiv preprint arXiv:2608.07208*, 2026.
- [19] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislav Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared Kaplan. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- [20] Jingwei Ni, Ekaterina Fadeeva, Tianyi Wu, Mubashara Akhtar, Jiaheng Zhang, Elliott Ash, Markus Leippold, Timothy Baldwin, See-Kiong Ng, Artem Shelmanov, and Mrinmaya Sachan. ReProbe: Efficient test-time scaling of multi-step reasoning by probing internal states of large language models. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2026.
- [21] Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Kotek, and Yonatan Belinkov. LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In *International Conference on Learning Representations (ICLR)*, 2025.
- [22] Nikunj Saunshi, Nishanth Dikkala, Zhiyuan Li, Sanjiv Kumar, and Sashank J. Reddi. Reasoning with latent thoughts: On the power of looped transformers. In *International Conference on Learning Representations (ICLR)*, 2025.
- [23] Tal Schuster, Adam Fisch, Jai Gupta, Mostafa Dehghani, Dara Bahri, Vinh Q. Tran, Yi Tay, and Donald Metzler. Confident adaptive language modeling. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [24] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.